

Elastic Metaspaces (Upstream Sales Pitch)

Thomas Stüfe, SAP
March 2020

PUBLIC

Agenda

- Motivation
- Basics
- Current implementation
- New implementation proposal

Motivation


- Reduce Memory Footprint
- Keep Metaspace coding maintainable

Basics

Metadata lifecycle

- Metadata are typically allocated when classes are loaded
- Accumulated metadata is freed after class loader has been collected
 - Bulk free scenario -> No need to track individual allocations -> arena style allocation
- Exceptions: premature deallocation possible but atypical

Metadata allocation / deallocation histogram (blocks/word size)



- Taken during a wildfly startup (standalone, no apps running)

- Small block allocations dominate. Only about 1% of allocations larger than 40 words.

- Deallocation: about 1:1000, similar curve.

Why bother?

Why not use a general purpose allocator (malloc, dlmalloc, boost etc?)

We think we can do better:

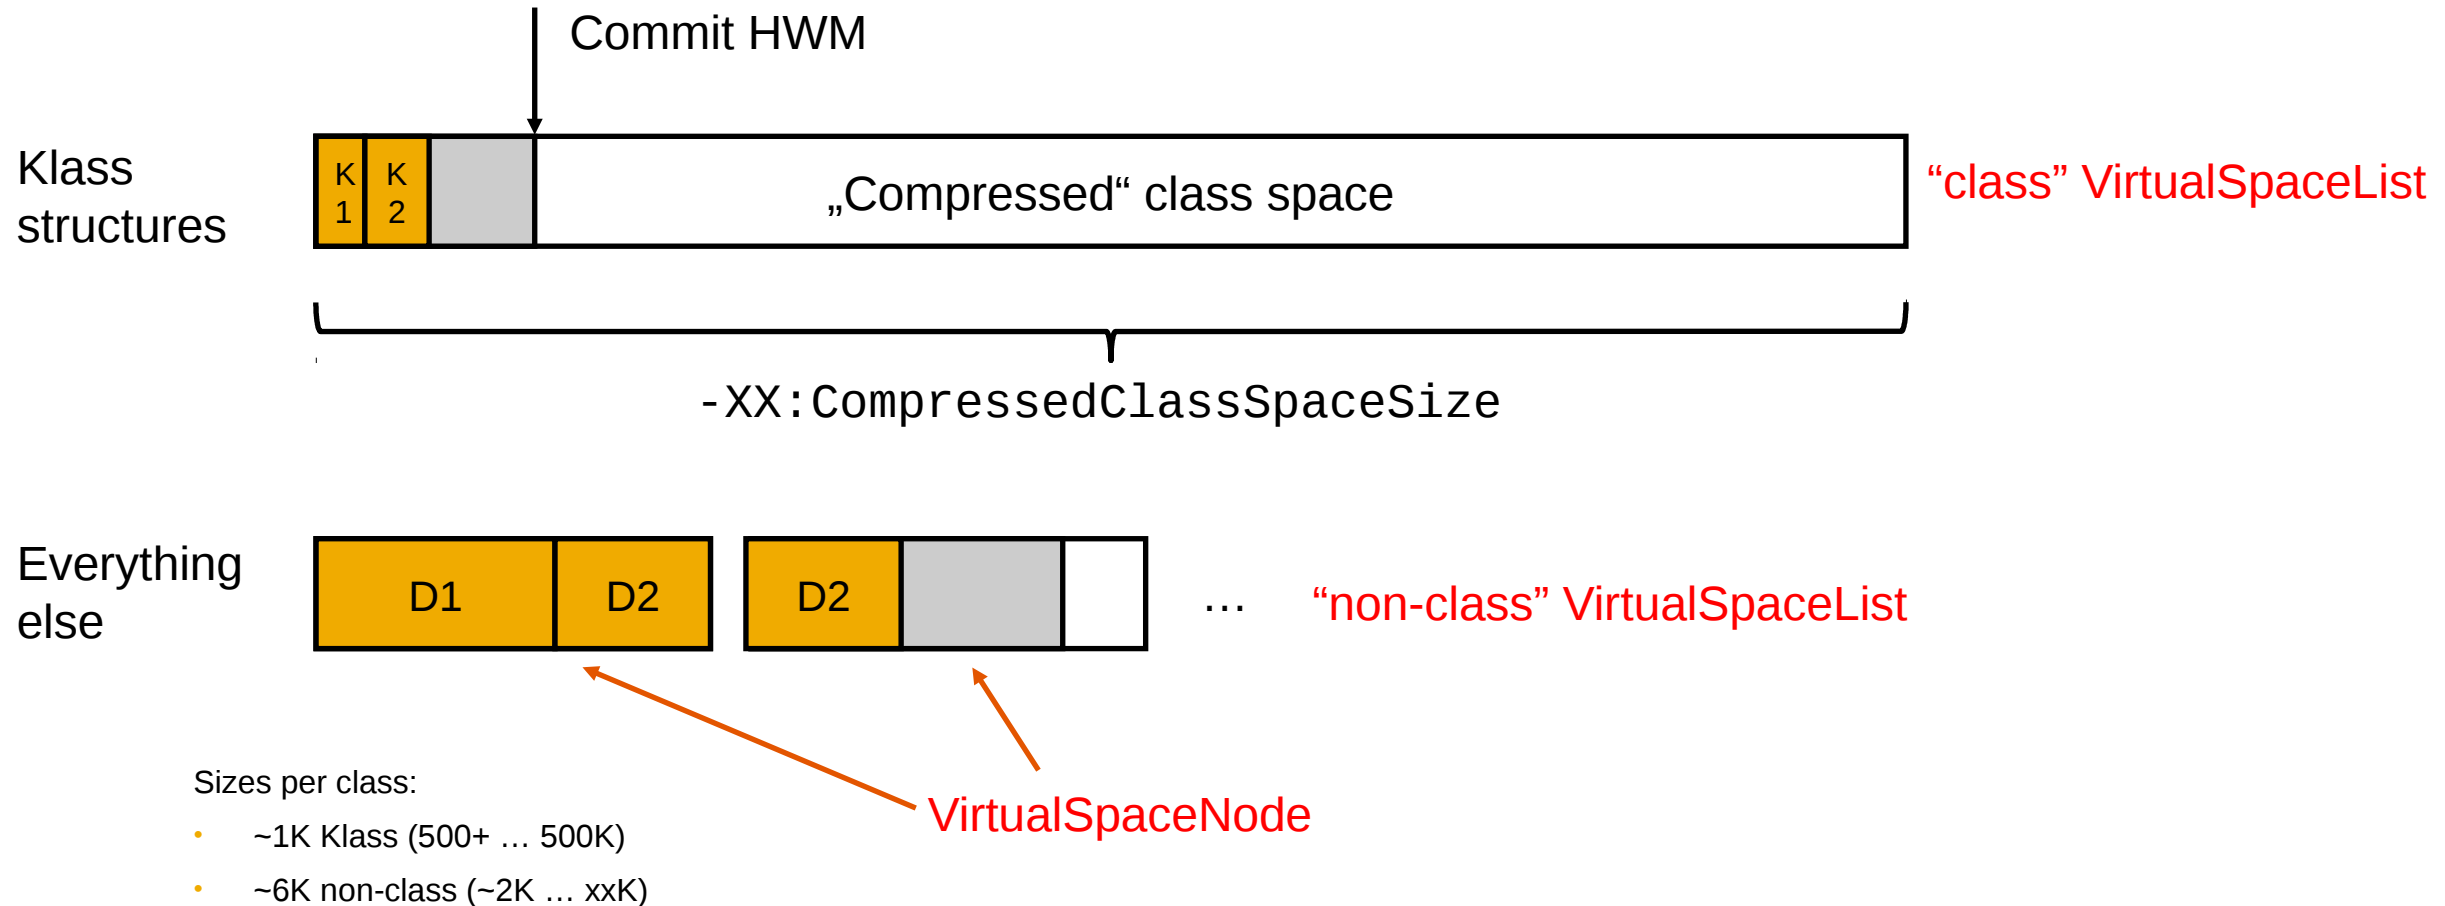
- Arena style allocation is fast and memory efficient.
- We know the size distribution of typical allocations
- malloc in particular would not work anyway:
 - CompressedClassSpace
 - Platform specific limitations (e.g. sbrk hits java heap)

Current implementation

Metaspace, very simplified, on one slide


- Lowest level: a series of memory regions, mmap'd, grows on demand, committed on demand.
- Class loaders allocate largish chunks from those regions.
- Metadata memory is allocated from those chunks via pointer bump.
- When a loader dies, its chunks are returned to a global freelist and may be reused.

Lowest level: VirtualSpaceList and VirtualSpaceNode



Current implementation

(much simplified)



- Loader owns a chunk of memory.
- Allocates from it via pointer bump.
 - Remember: we do not need to track individual allocations for freeing.

Current implementation (2)


(much simplified)



- If chunk is used up, Loader acquires a new one from the metaspace allocator.
- Retired chunks are kept in list
- Leftover space is kept for later reuse

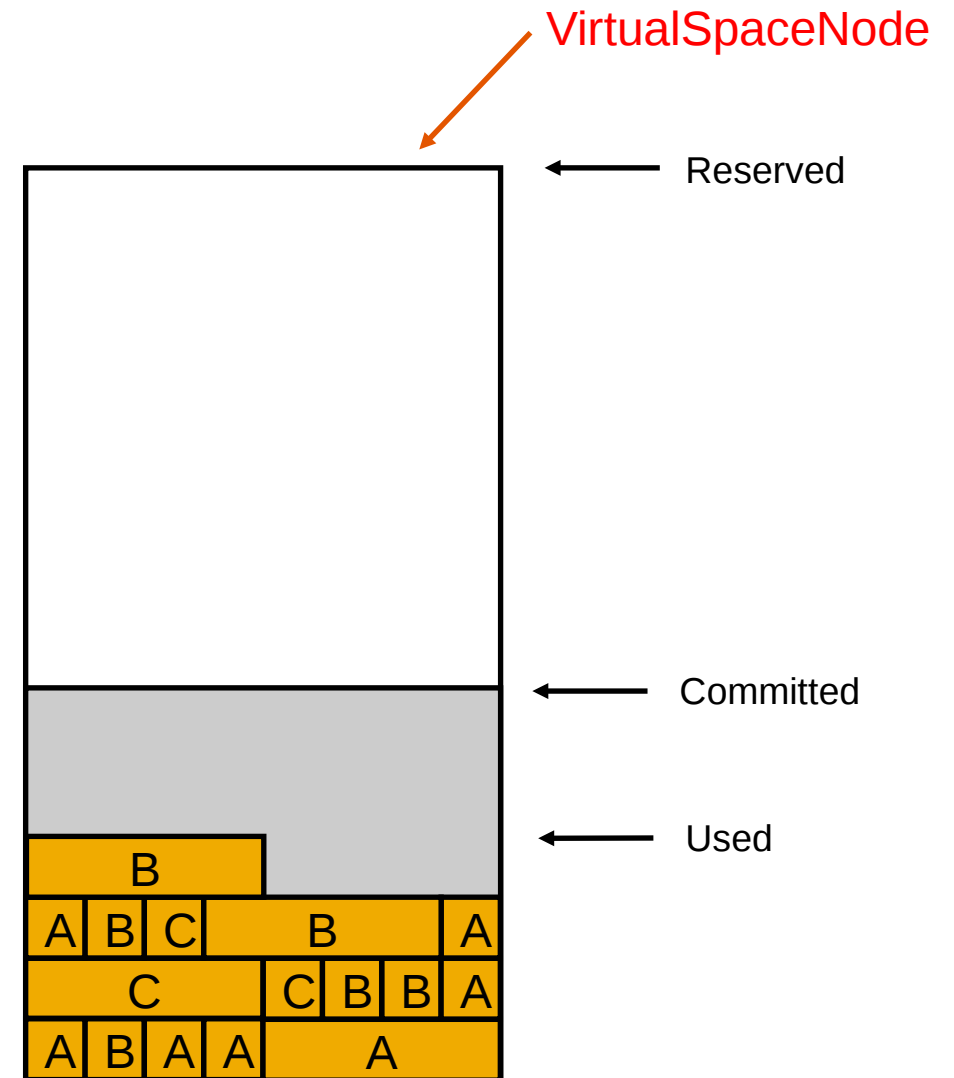
Current implementation (3)

(much simplified)




Chunks are carved from **VirtualSpaceNode** as they are allocated.

VirtualSpaceNode is committed on demand, never gets uncommitted.

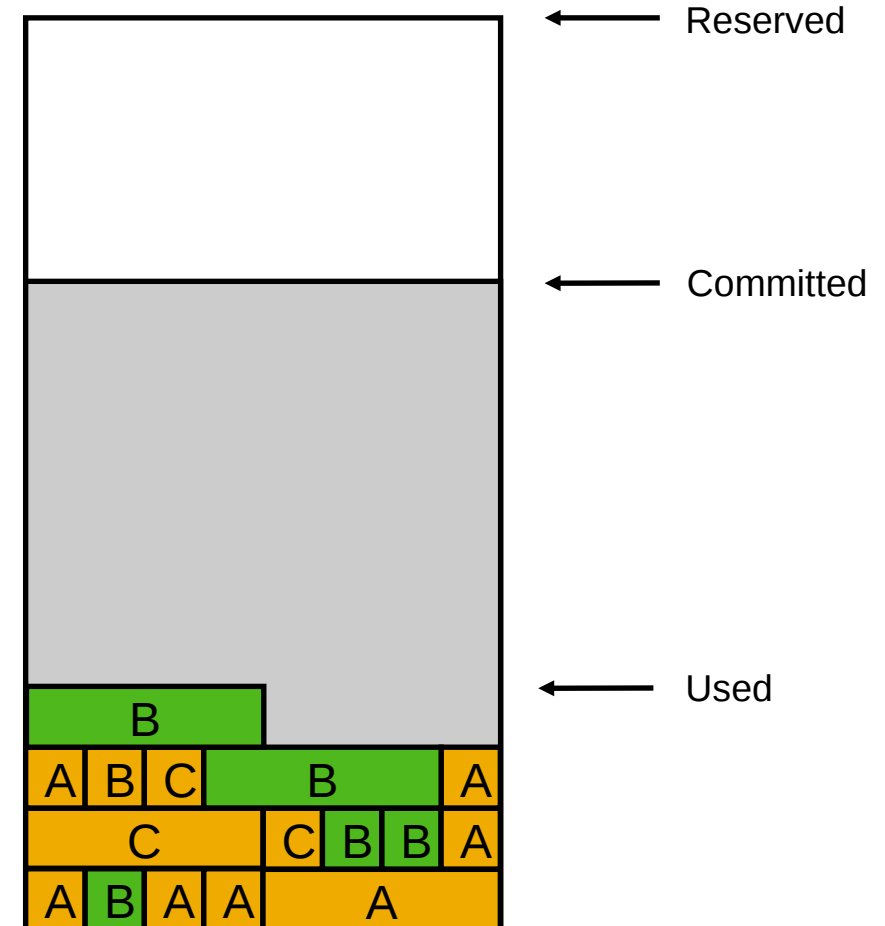


Current implementation (4)

(much simplified)

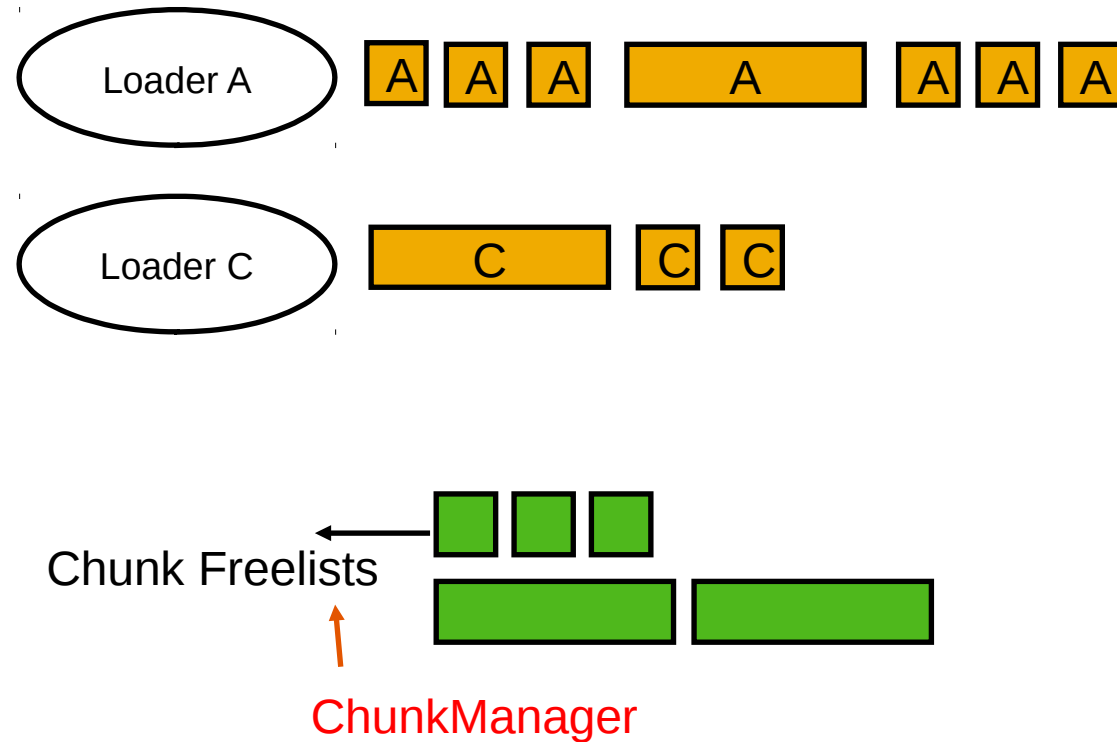


When a loader dies, its chunks are marked as free...

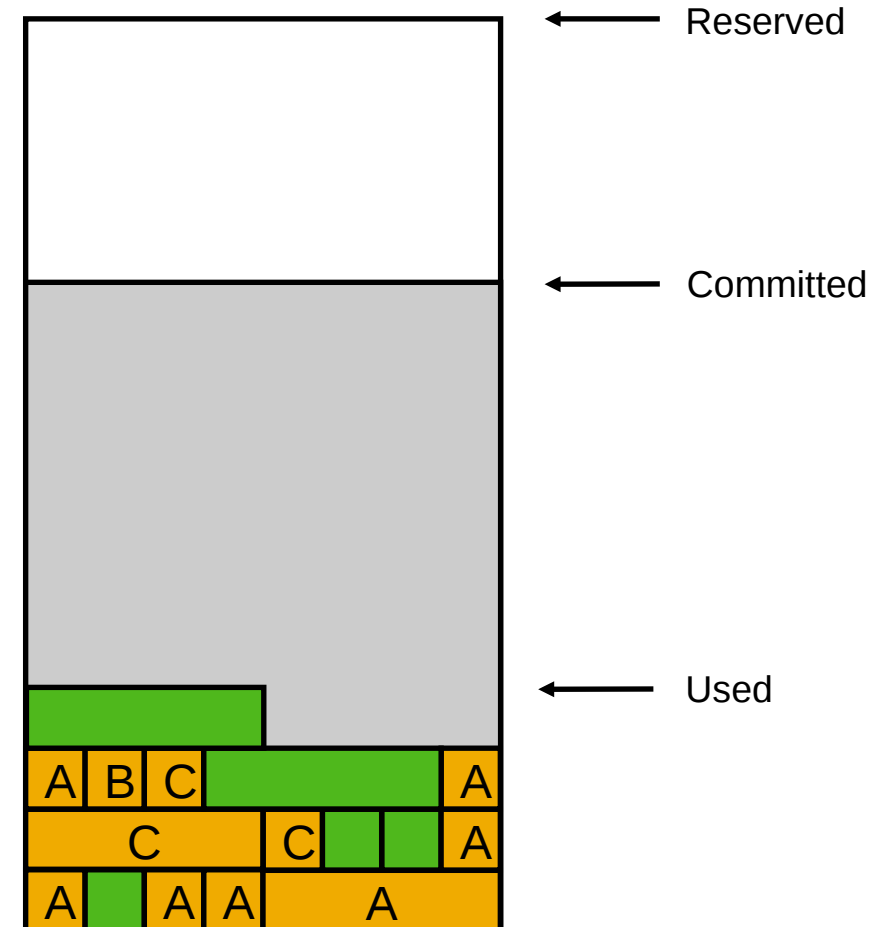


Current implementation (5)

(very much simplified)




...and added to global freelists (**ChunkManager**), sorted by size. **VirtualSpaceNode** is potentially unmapped.



Deallocations

- Premature release of metadata
 - 1 Class redefinitions, Profile Counters, partly loaded metadata on Class loading errors ..
 - 2 But also: remains of retired chunks (usually larger blocks)
- Typically very rare (~ 1:1000 dealloc:alloc). But may happen more often in pathological cases
 - Instrumentation?
- Deallocated metadata are still owned by class loader
- We attempt to reuse them for follow up metadata allocation, to varying degrees of success
- Deallocation histogram: similar to allocation but higher number of larger chunks due to (2)


Deallocation




Deallocation: BlockFreeList

- Two parts
 - „SmallBlocks“ - a ordered vector of linked lists of free metadata, for small sizes (up to 13 blocks)
 - $O(1)$ insert/retrieval
 - „BlockTreeDictionary“ - a BST for larger blocks
 - BinaryTreeDictionary
 - Unbalanced
- Whats not optimal:
 - Smallblocks does not search for larger blocks
 - Dictionary: too large blocks are retrieved and inserted back unnecessarily
 - We are hesitant to query the BTD
 - Complicated code

Classes




Per class loader



Global


..with CompressedClassPointers



Class metadata

Per class loader


“non-class” metadata



Class VSlist (note: really just one element).

Non-Class VSlist

Global



Class ChunkManager

Non-Class ChunkManager

Chunk sizes

- Class space: 1K (“special”), 2K (“small”), 32K (“medium”)
- Non-class space: 1K, 4K, 64K
- And humongous chunks: larger than medium and variable sized

Metachunk coalescation

- The „chunk size choking problem“
- JDK-8198423
- Since then chunks can be merged and split, within limits
 - 4x1K chunks -> 4K
 - 16x4K chunks -> 64K
- Basically the whole thing is now a weird buddy allocator
 - A bit inefficient due to the odd chunk geometry
 - But it solved the problem
 - But I was afraid to touch too much, so the whole patch is one gigantic band aid
 - Ugly and difficult to maintain :-(

Concern: „Micro Loaders“


- Some loaders / CLDs only ever load one class:
 - Anonymous classes
 - Reflection delegators

- Not optimal.
 - In class space, only one InstanceClass is allocated, but needs a whole chunk
 - In non-class space, ~10-20 allocations

Main waste areas

- Freelists can get huge.
 - We have seen used:free ratios of 1:3 and worse
 - =>Metaspace is not really elastic.
- Intra-chunk waste
 - At some point loader typically stops loading classes; remaining chunk space (and deallocated space waiting for reuse) is wasted
 - Worse with micro loaders, if there are a lot

Huge freelists: Committed vs used space, after class unloading



Huge Freelists (jcmd VM.metaspaces output)

```
jcmd 27265 VM.metaspaces
```

```
27265:
```

```
...
```

```
Waste (percentages refer to total committed size 373,48 MB):
```

| | |
|---------------------------------|--------------------------------|
| Committed unused: | 280,00 KB (<1%) |
| Waste in chunks in use: | 2,45 KB (<1%) |
| Free in chunks in use: | 6,34 MB (2%) |
| Overhead in chunks in use: | 186,75 KB (<1%) |
| In free chunks: | 269,56 MB (72%) |
| Deallocated from chunks in use: | 998,98 KB (<1%) (1763 blocks) |
| -total-: | 277,33 MB (74%) |

Eat up! - Intrachunk waste from idle loaders


- Most loaders stop loading at some point. Remainder space in current chunk as well as deallocated blocks remain unused, effectively wasted.
- How large a chunk do we give to a loader?
 - Too small: high fragmentation and contention on central allocator parts
 - Too large: wasted space
 - → we try to guess future loading behavior. We may guess wrong and suffer.
- We have tuned these areas a lot (see e.g. Zhengyu Yu's work on JDK-8190729 and JDK-8191924) but the solutions are far from optimal

Reimplementation

Basic idea

- Chunks in freelists can be uncommitted
- Delay committing chunks until they are actually used
 - Partly commit them piece wise (like a thread stack)
 - Removes the penalty of handing out large chunks to class loaders

Commit granules



Current chunk allocation scheme unsuited for uncommitting chunks

- Odd chunk geometry
 - Difficult to merge and split
 - High fragmentation
 - Complex code
- Chunk headers are a problem


Pow 2 based buddy allocator for chunks

- We need a better, cleaner chunk geometry scheme.
- Power 2 based buddy allocation is a nice fit.
- Chunks sized from 1K ... 4M in pow2 steps (13 sizes in total).
- Dead simple to split and merge.
- Low external defragmentation -> Leads to larger free contiguous areas.
- Standard algorithm widely known

Pow 2 based buddy allocator for chunks


```
// Each chunk has a level; the level corresponds to its position in the tree
// and describes its size.
//
// The largest chunks are called root chunks, of 4MB in size, and have level 0.
// From there on it goes:
//
// size      level
// 4MB       0
// 2MB       1
// 1MB       2
// 512K      3
// 256K      4
// 128K      5
// 64K       6
// 32K       7
// 16K       8
// 8K        9
// 4K        10
// 2K        11
// 1K        12
```

Buddy allocator: Allocation



- Remove chunk from freelist
- Optionally split until desired size is reached
- Return result chunk; put splinter chunks back to freelist

Buddy allocator: Deallocation



- Mark chunk as free
- If buddy chunk is free and unsplit: remove from freelist and merge with chunk
 - Repeat until root chunk sized reached or until buddy is not free
- Return result chunk to free list


New VirtualSpaceNode

- Metaspace is now segmented into “root chunk areas”: 4MB sized, aligned to 4MB
- VirtualSpaceNode:
 - now only allocates root chunks → much simplified code.
 - Does not commit! But provides committing as service to upper layers.
 - Keeps a bitmap to keep track of committed/uncommitted granules.

New Chunk Manager

- ChunkManager keeps 13 freelists, one per chunk level
- Allocation flow:
 - SpaceManager to ChunkManager: give me chunk of level X
 - ChunkManager: have one in freelist(X) – ok.
 - Have none? Search upward in freelists(X+1,2...)
 - Found a larger chunk? Split it buddy style; return chunk of level X; put splinter chunks back into respective freelists
 - Found no larger chunk? Ask underlying VirtualSpaceNode for new root chunk. Proceed as described above.


Granules and chunks – putting it all together



A larger chunk can span multiple granules (1:n)

Multiple small chunks can cover a single granule (n:1)


Granules and chunks – putting it together



Free chunks spanning 1+ granules can be uncommitted

A chunk spanning >1 granules can be committed in parts, on demand


Larger chunks can be committed on demand




Removes penalty for guessing wrong in the “how large a chunk should I give him” guessing game at least for larger chunks (e.g. for the boot loader).

Chunk headers needed to go

Before



Now



New Deallocation handling („LeftOverManager“)

- Bin list (similar to SmallBlocks) + a newly written BST
 - New binlist covers more sizes (atm 32) and searches also upward
 - The new BST is similar to the old one, but much reduced in code size and much simpler.
 - New BST knows its largest node size.
 - Note: we do not need the binaryTreeDictionary anymore :)
- We now split blocks where it makes sense.
- Possible further improvement: make it an RB tree.

What else changed

- Got rid of humongous chunks :)
- Got rid of occupancy map
- Nice: Chunks can now grow in-place
 - Saves overhead and reduces intrachunk waste
- Code cleaner and more maintainable; better separation of concerns and testability.
- Lots of new gtests

More improvements possible

- Improve error analysis for overwrites in Metaspace
 - Simple: Buffer zones, canaries, disabling deallocation
 - Costly but possible: guard pages
- Better micro loader handling


Concern 1: keep number of virtual memory areas low

- (Linux):
 - Higher commit/uncommit fragmentation results in higher number of VMAs
 - Kernel keeps vma structures in list and rb tree
 - Too many of them may affect vma lookup
 - And we may hit process limits
- So: keep an eye on commit granularity
- Solution: commit granule size is adjustable. Larger granules → lower fragmentation at the cost of lower memory returns.


Concern 2: uncommit speed

- Matters only to GCs which have no concurrent class unloading
- (Linux):
 - Page table has to be unrolled & deallocated. How expensive this is depends on population of uncommitted area: how many pages had been committed before, and their size.
 - Hence indirectly on size of uncommit region
 - And also on number of vma in committed region, although in our case it should always be one.
- Currently I see no problems in practice, but a fall-back plan would be to uncommit concurrently. Not that complicated if we keep current locking scheme


Result: Committed vs used, Stock JDK14




Result: Committed vs used, Patched JDK14



Result: committed Metaspace, Stock vs Patched VM



Result: RSS, Stock vs Patched VM



Modest decrease in consumption beyond class unloading

- Wildfly standalone after startup: 61m->54m, -7m, (11%)
- Eclipse CDT, hotspot project after C++ indexing: 138m->129m, -9m (12%)
- jruby helloworld.rb (invokedynamic, compile=FORCE): 41m->38m, -3m, (1.2%)

Performance costs

- jbb2015 did not show regressions
- Micro benchmarks doing mass class loading and unloading show atm 1-2% decrease in performance. I am working on it.

How do we go from here?

- Patch is stable. Needs more tests and smaller fixes but it works.
- Patch lives in jdk/sandbox repository, branch "stuefe-new-metaspace-branch"
 - <http://hg.openjdk.java.net/jdk/sandbox/>
- JEP exists in Draft state ("Elastic Metaspace": <https://openjdk.java.net/jeps/8221173>)
- JDK15?
- A good candidate for backporting
 - Would make a lot of sense in 11/8
 - Large patch but Metaspace is quite isolated. Should not be too much of a hassle.

Thank you.

Contact information:

Thomas Stüfe

[@tstuefe](https://www.instagram.com/tstuefe)

thomas.stuefe@sap.com

[stuefe.de](https://www.stuefe.de)